

# Expression templates for primal value taping in the reverse mode of AD

Max Sagebaum, Tim Albring and Nicolas R. Gauger

AG Scientific Computing  
TU Kaiserslautern

AD2016 - 7th International Conference on Algorithmic Differentiation

# Overview

- Motivation: Jacobi taping vs. primal value taping
- Expression templates for primal value taping
- Handling of constant and passive types
- Memory and time analysis

# CoDiPack introduction

Active type structure:

```
struct RealReverse {  
    double p; // primal value  
    int i;    // index for the adjoint  
}
```

Index management

- Linear index manager (e.g. RealReverse)
  - Statements and adjoint variables tightly coupled
  - $i = ++\text{tape.globalStatementIndex};$
  - Advantage: LHS indices can be computed (no storing of LHS index)
  - Drawback: Big adjoint vector
- Reuse index manager (e.g. RealReverseIndex)
  - No coupling between statements and adjoint variables
  - $\text{tape.indexManager.assignIndex}(i);$
  - Advantage: Small adjoint vector
  - Drawback: Additional 4 bytes per statement (index needs to be stored)

# CoDiPack introduction

## Tape implementation

- Simple (e.g. RealReverseUnchecked)
  - No bounds check performed
  - Memory needs to be preallocated
  - Faster evaluation
- Chunk (e.g. RealReverse)
  - Memory allocated on the fly
  - Slower due to bounds check

# Motivation: Jacobi taping vs. primal value taping

Elemental operator

$$\phi : \mathbb{R}^n \rightarrow \mathbb{R}, n \in \mathbb{N}$$

Reverse update

$$\bar{x}_i += \frac{\partial \phi}{\partial x_i}(x_i)^T \bar{y}, \quad \forall i = 1 \dots n$$

$$\bar{y} = 0$$

# Motivation: Jacobi taping vs. primal value taping

Elemental operator

$$\phi : \mathbb{R}^n \rightarrow \mathbb{R}, n \in \mathbb{N}$$

Reverse update

$$\bar{x}_i += \frac{\partial \phi}{\partial x_i}(x_i)^T \bar{y}, \quad \forall i = 1 \dots n$$

$$\bar{y} = 0$$

Jacobi taping:

- Store  $\frac{\partial \phi}{\partial x_i}$ , indices for  $x_i$  and  $n$
- Bytes per statement:  $n * 8 + n * 4 + 1 = 12 * n + 1$
- Assumption:  $n < 256 \Rightarrow$  Each elemental operation less than 256 arguments

# Motivation: Jacobi taping vs. primal value taping

Elemental operator

$$\phi : \mathbb{R}^n \rightarrow \mathbb{R}, n \in \mathbb{N}$$

Reverse update

$$\bar{x}_i += \frac{\partial \phi}{\partial x_i}(x_i)^T \bar{y}, \quad \forall i = 1 \dots n$$

$$\bar{y} = 0$$

Jacobi taping:

- Store  $\frac{\partial \phi}{\partial x_i}$ , indices for  $x_i$  and  $n$
- Bytes per statement:  $n * 8 + n * 4 + 1 = 12 * n + 1$
- Assumption:  $n < 256 \Rightarrow$  Each elemental operation less than 256 arguments

Primal value taping:

- Store  $y$ , indices for  $x_i$  and handle to function
- Bytes per statement:  $8 + n * 4 + 8 = 4 * n + 16$

# Expression templates for primal value taping

The equation

$$w = ((a + b) * (c - d))^2$$



# Expression templates for primal value taping

The equation

$$w = ((a + b) * (c - d))^2$$

is represented by the structure

POW<MULT<ADD<ActiveReal, ActiveReal>, SUB<ActiveReal, ActiveReal>>>

# Expression templates for primal value taping


The equation

$$w = ((a + b) * (c - d))^2$$

is represented by the structure

POW<MULT<ADD<ActiveReal, ActiveReal>, SUB<ActiveReal, ActiveReal>>>

The reverse AD graph:



# Expression templates for primal value taping


The equation

$$w = ((a + b) * (c - d))^2$$

is represented by the structure

POW<MULT<ADD<ActiveReal, ActiveReal>, SUB<ActiveReal, ActiveReal>>>

The reverse AD graph:



The required code:

```
add = a + b; sub = c - b; mul = add * sub;
w = pow(mul, 2.0);

w_b = tape.getAdjoint(w_i);
mul_b += 2 * mul * w_b;
sub_b += add * mul_b;
add_b += sub * mul_b;
c_b += sub_b;
d_b += -sub_b;
a_b += add_b;
b_b += add_b;
tape.updateAdjoint(a_b, a_i);
tape.updateAdjoint(b_b, b_i);
tape.updateAdjoint(c_b, c_i);
tape.updateAdjoint(d_b, d_i);
```

## Expression templates for primal value taping

Which actions do we need?

## Expression templates for primal value taping

Which actions do we need?

- 1 Gather all indices
- 2 Evaluate the adjoint
- 3 Store a function handle

? .. Probably many more

## Expression templates for primal value taping - 1. Gather

- Add per value action to expression interface

```
template<typename A>
struct Expression {
    ...
    template<typename Data, typename Func>
    inline void valueAction(Data data, Func func) const;
}
```

## Expression templates for primal value taping - 1. Gather

- Add per value action to expression interface

```
template<typename A>
struct Expression {
    ...
    template<typename Data, typename Func>
    inline void valueAction(Data data, Func func) const;
}
```

```
template<typename A, typename B>
struct BinaryOperator : public Expression<BinaryOperator<A, B> > {
    template<typename Data, typename Func>
    inline void valueAction(Data data, Func func) const {
        a_.valueAction(data, func);
        b_.valueAction(data, func);
    }
}
```

## Expression templates for primal value taping - 1. Gather

- Add per value action to expression interface

```
template<typename A>
struct Expression {
    ...
    template<typename Data, typename Func>
    inline void valueAction(Data data, Func func) const;
}
```

```
template<typename A, typename B>
struct BinaryOperator : public Expression<BinaryOperator<A, B> > {
    template<typename Data, typename Func>
    inline void valueAction(Data data, Func func) const {
        a_.valueAction(data, func);
        b_.valueAction(data, func);
    }
}
```

```
struct RealReverse: public Expression<RealReverse> {
    template<typename Data, typename Func>
    inline void valueAction(Data data, Func func) const {
        CALL_MEMBER_FN(globalTape, func)(data, this->p, this->i);
    }
}
```




# Expression templates for primal value taping - 1. Gather

## ■ Use the method in the tape

```
struct PrimalValueTape {  
    ...  
    template<typename Expr>  
    inline void store(double& lhsValue, int& lhsIndex, const Expr& rhs) {  
        indexVector.reserveItems(ExpressionTraits<Expr>::maxActiveVariables);  
        rhs.valueAction(NULL, &PrimalValueTape::pushIndices);  
    }  
    ...  
}  
  
inline void pushIndices(void* data, const double& value, const int& index) {  
    indexVector.setDataAndMove(pushIndex);  
}  
}
```

## Expression templates for primal value taping - 2. Evaluate



## Expression templates for primal value taping - 2. Evaluate

```
// BinaryOperator
template<size_t offset>
static inline double getValue(const int* indices, const Real* primalValues) {
    const double aPrimal = A::template getValue<offset>(indices, primalValues);
    const double bPrimal = B::template getValue
        <offset + ExpressionTraits<A>::maxActiveVariables>(indices, primalValues);
    return PRIMAL_CALL(aPrimal, bPrimal);
}
```

## Expression templates for primal value taping - 2. Evaluate

```
// BinaryOperator
template<size_t offset>
static inline double getValue(const int* indices, const Real* primalValues) {
    const double aPrimal = A::template getValue<offset>(indices, primalValues);
    const double bPrimal = B::template getValue
        <offset + ExpressionTraits<A>::maxActiveVariables>(indices, primalValues);
    return PRIMAL_CALL(aPrimal, bPrimal);
}
```

```
// RealReverse
template<size_t offset>
static inline double getValue(const int* indices, const Real* primalValues) {
    return primalValues[indices[offset]];
}
```

## Expression templates for primal value taping - 2. Evaluate

```
//BinaryOperator
template<size_t offset>
static inline void evalAdjoint(const double& seed, const int* indices,
                               const Real* primalValues, Real* adjointValues) {
    const double aPrimal = A::template getValue<offset>(indices, primalValues);
    const double bPrimal = B::template getValue
        <offset + ExpressionTraits<A>::maxActiveVariables>(indices, primalValues);
    const double resPrimal = PRIMAL_CALL(aPrimal, bPrimal);

    const double aJac = GRADIENT_FUNC_A(aPrimal, bPrimal, resPrimal) * seed;
    const double bJac = GRADIENT_FUNC_B(aPrimal, bPrimal, resPrimal) * seed;
    A::template evalAdjoint<offset>(aJac, indices, primalValues, adjointValues);
    B::template evalAdjoint<offset + ExpressionTraits<A>::maxActiveVariables>
        (bJac, indices, primalValues, adjointValues);
}
```

## Expression templates for primal value taping - 2. Evaluate

```
//BinaryOperator
template<size_t offset>
static inline void evalAdjoint(const double& seed, const int* indices,
                               const Real* primalValues, Real* adjointValues) {
    const double aPrimal = A::template getValue<offset>(indices, primalValues);
    const double bPrimal = B::template getValue
        <offset + ExpressionTraits<A>::maxActiveVariables>(indices, primalValues);
    const double resPrimal = PRIMAL_CALL(aPrimal, bPrimal);

    const double aJac = GRADIENT_FUNC_A(aPrimal, bPrimal, resPrimal) * seed;
    const double bJac = GRADIENT_FUNC_B(aPrimal, bPrimal, resPrimal) * seed;
    A::template evalAdjoint<offset>(aJac, indices, primalValues, adjointValues);
    B::template evalAdjoint<offset + ExpressionTraits<A>::maxActiveVariables>
        (bJac, indices, primalValues, adjointValues);
}
```

```
//RealReverse
template<size_t offset>
static inline void evalAdjoint(const double& seed, const int* indices,
                               const Real* primalValues, Real* adjointValues) {
    adjointValues[indices[offset]] += seed;
}
```

## Expression templates for primal value taping - 3. Store

### ■ Required information:

- Function pointer (8 byte)
- Number of arguments (1 byte)
- $\Rightarrow$  Static information does not need to be stored in the tape

$\Rightarrow$  Use static handles and store references to them:

```
// the handle
struct ExpressionHandle {
    const StatementFuncPointer adjointFunc;
    const size_t maxActiveVariables;
};

// the store
template<typename Expr>
struct ExpressionHandleStore {
    static const ExpressionHandle handle;
};

// the instantiation
template<typename Expr>
const ExpressionHandle ExpressionHandleStore<Expr>::
    handle(Expr::template evalAdjoint<0>, ExpressionTraits<Expr>::maxActiveVariables);
```

# Expression templates for primal value taping

Complete store function:

```
// PrimalValueTape
template<typename Expr>
inline void store(double& lhsValue, int& lhsIndex, const Expr& rhs) {
    indexVector.reserveItems(ExpressionTraits<Expr>::maxActiveVariables);
    rhs.valueAction(NULL, &PrimalValueTape::pushIndices);

    stmtVector.reserveItems(1);
    stmtVector.setDataAndMove(&ExpressionHandleStore<Rhs>::handle);

    lhsIndex = ++this->globalStatementIndex;
    lhsValue = rhs.getValue();

    checkPrimalsSize();
    this->primals[lhsIndex] = lhsValue;
}
```



# Expression templates for primal value taping

How to evaluate in the reverse sweep:

```
// PrimalValueTape
inline void evaluateStack(const size_t& startAdjPos, const size_t& endAdjPos,
                        size_t& stmtPos, Handle* &statements,
                        size_t& indexPos, IndexType* &indices) {
    size_t adjPos = startAdjPos;

    while(adjPos > endAdjPos) {
        const GradientValue& adj = adjoints[adjPos];
        --adjPos;
        --stmtPos;

        indexPos -= statements[stmtPos].maxActiveVariables;
        statements[stmtPos].adjointFunc(adj, &indices[indexPos], primalVector, adjoints);
    }
}
```

# Expression templates for primal value taping

Summary expression templates for primal value taping:

- Use general actions to store the indices
- Use static handles to store the **static** data for each statement
  - Function pointer
  - Number of variables
- Optimized implementation w.r.t. the compiler:
  - Offset into arrays as template parameter
  - Inlining of recursive calls

# Expression templates for primal value taping

Summary expression templates for primal value taping:

- Use general actions to store the indices
- Use static handles to store the `static` data for each statement
  - Function pointer
  - Number of variables
- Optimized implementation w.r.t. the compiler:
  - Offset into arrays as template parameter
  - Inlining of recursive calls

What is missing?

# Expression templates for primal value taping

Summary expression templates for primal value taping:

- Use general actions to store the indices
- Use static handles to store the `static` data for each statement
  - Function pointer
  - Number of variables
- Optimized implementation w.r.t. the compiler:
  - Offset into arrays as template parameter
  - Inlining of recursive calls

What is missing?

- Constant values
- Passive values

# Constant value handling

Constant value example:

$$w = \sin(2.0 * a)$$

Required changes:

- Add push passive action

```
inline void pushPassive(void*, const double& value) {  
    constantValueVector.setDataAndMove(value);  
}
```

- Add constant value count to handle

```
constantPos -= exprHandle->maxConstantVariables;
```

- Add constant offset to evaluation function

```
template<size_t offset, size_t constantOffset>  
static inline void evalAdjoint(  
    const double& seed, const int* indices, const double* constantValues,  
    const double* primalValues, double* adjointValues);
```

# Passive value handling

Passive value example:

```
RealReverse two = 2.0; // has index 0  
RealReverse three = 3.0; // has index 0  
RealReverse w = sin(two * a) / three; // which value should be stored in index 0
```

Possible approaches:

# Passive value handling

Passive value example:

```
RealReverse two = 2.0; // has index 0  
RealReverse three = 3.0; // has index 0  
RealReverse w = sin(two * a) / three; // which value should be stored in index 0
```

Possible approaches:

- 1 Make variable active in statement
  - Each use creates an additional 24 bytes

# Passive value handling

Passive value example:

```
RealReverse two = 2.0; // has index 0  
RealReverse three = 3.0; // has index 0  
RealReverse w = sin(two * a) / three; // which value should be stored in index 0
```

Possible approaches:

- 1 Make variable active in statement
  - Each use creates an additional 24 bytes
- 2 Activate variable on first use
  - Only possible with reuse index manager
  - First use creates an additional 24 bytes
  - Increases active section of program  
⇒ more statements recorded



# Passive value handling

Passive value example:

```
RealReverse two = 2.0; // has index 0
RealReverse three = 3.0; // has index 0
RealReverse w = sin(two * a) / three; // which value should be stored in index 0
```

Possible approaches:

- 1 Make variable active in statement
  - Each use creates an additional 24 bytes
- 2 Activate variable on first use
  - Only possible with reuse index manager
  - First use creates an additional 24 bytes
  - Increases active section of program  
⇒ more statements recorded
- 3 Store values in passive vector
  - Global use of an additional  $256 * 24$  bytes
  - Each use create an additional 8 bytes
  - Each statement used an additional 1 byte

## Passive value handling - Option 3

Required changes:

- Store passive value with indices

```
inline void pushIndices(int* passiveVarCount, const double& value, const int& index) {  
    IndexType pushIndex = index;  
    if(0 == pushIndex) {  
        *passiveVarCount += 1;  
        pushIndex = *passiveVarCount;  
        constantValueVector.setDataAndMove(value);  
    }  
  
    indexVector.setDataAndMove(pushIndex);  
}
```

- Add number of passives to statement

```
stmtVector.setDataAndMove(handle, (uint8_t)passiveVariableCount);
```

- Update primal values in reverse sweep

```
constantPos -= passiveActives;  
for(StatementInt i = 0; i < passiveActives; ++i) {  
    primalVector[i + 1] = constants[constantPos + i];  
}
```

## Example and memory consumption

- 2D coupled Burgers equation on 601x601 grid with 32 time steps


$$u_t + uu_x + vu_y = \frac{1}{R}(u_{xx} + u_{yy}) \quad (1)$$

$$v_t + uv_x + vv_y = \frac{1}{R}(v_{xx} + v_{yy}) \quad (2)$$


- Evaluated on one node of the Elwetritsch cluster with two Intel E5-2670 cpu's (16 cores)
- *Single* test case: Run only one process
  - Full memory bandwidth available
- *Multi* test case: Run 16 times the same problem
  - Memory bandwidth is limited

	Jacobi	Jacobi index	Primal	Primal index
Memory in MB	4830	4496	4030	3853
Reduction	0 %	7 %	17 %	20 %


# Time measurements - Recording time




## Time measurements - Recording time




# Time measurements - Recording time



# Time measurements - Recording time




## Time measurements - Reverse evaluation time





## Time measurements - Reverse evaluation time



# Conclusion & Guidelines & Outlook

## Conclusion:

- Fast implementation of primal value taping with expression templates
- Saves memory w.r.t. Jacobi taping
- Passive values need to be looked at

# Conclusion & Guidelines & Outlook

## Conclusion:

- Fast implementation of primal value taping with expression templates
- Saves memory w.r.t. Jacobi taping
- Passive values need to be looked at

## Guidelines:

- Don't use constants e.g. 4.0
- Use active constants e.g.

```
RealReverse FOUR = 4.0;  
tape.registerInput(FOUR);
```

- Legacy codes with a lot of passive values may prefer Jacobi taping

# Conclusion & Guidelines & Outlook

## Conclusion:

- Fast implementation of primal value taping with expression templates
- Saves memory w.r.t. Jacobi taping
- Passive values need to be looked at

## Guidelines:

- Don't use constants e.g. 4.0
- Use active constants e.g.

```
RealReverse FOUR = 4.0;  
tape.registerInput(FOUR);
```

- Legacy codes with a lot of passive values may prefer Jacobi taping

## Outlook:

- Investigate inline problems
- Implement different passive value handling for index tapes
- Improve memory per statement (from 16 byte to 10 byte)

# CoDiPack Release 1.3

## CoDiPack release 1.3

- Primal value tape implementation
  - RealReversePrimal, RealReversePrimalIndex
  - RealReversePrimalVec, RealReversePrimalIndexVec
  - RealReversePrimalUnchecked, RealReversePrimalIndexUnchecked
- *numeric\_limits* specialization
- Added *erf* and *erfc*
- Newsletter: [codi-info@uni-kl.de](mailto:codi-info@uni-kl.de)
- Contact: [codi@scicomp.uni-kl.de](mailto:codi@scicomp.uni-kl.de)

Thank you for your attention!